



Star

893

Fork

413

Watch

97

Follow @collabnix

193

Cheatsheet - Kubectl

Docker - Beginners | Intermediate | Advanced

[View on GitHub](#)[Join Slack](#)

Cheatsheet - Kubectl

Kubectl is a command line interface for running commands against Kubernetes clusters.

Installing

The kubectl version has to be within one minor version difference of the Kubernetes cluster. For example, a v1.2 client should work with v1.1, v1.2, and v1.3 master.

Kubectl can be installed on Ubuntu, Debian, CentOS, RedHat operating systems.

Ubuntu / Debian

```
sudo apt-get update && sudo apt-get install -y apt-transport-https
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee -a /etc/apt/sources.list.d/kubernetes.list
sudo apt-get update
sudo apt-get install -y kubectl
```

CentOS / RedHat

```
cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.
EOF
yum install -y kubectl
```

For further information about kubectl installation method, please refer to [the Kubernetes documentation](#).

Completion

To easy manage the Kubernetes resources thanks to the command line Kubectl, the shell completion can be added to the shell profile to easily navigate in command line.

```
# Installing bash completion on macOS using homebrew
## If running Bash 3.2 included with macOS
brew install bash-completion
## or, if running Bash 4.1+
brew install bash-completion@2
## If kubectl is installed via homebrew, this should start working immediately.
## If you've installed via other means, you may need add the completion to your comp
kubectl completion bash > $(brew --prefix)/etc/bash_completion.d/kubectl

# Installing bash completion on Linux
## Load the kubectl completion code for bash into the current shell
source <(kubectl completion bash)
## Write bash completion code to a file and source if from .bash_profile
kubectl completion bash > ~/.kube/completion.bash.inc
printf "
# Kubectl shell completion
source '$HOME/.kube/completion.bash.inc'
" >> $HOME/.bash_profile
source $HOME/.bash_profile

# Load the kubectl completion code for zsh[1] into the current shell
source <(kubectl completion zsh)
```

```
# Set the kubectl completion code for zsh[1] to autoload on startup
kubectl completion zsh > "${fpath[1]}/_kubectl"
```

Syntax

Kubectl is a powerful tool to manage each object on a Kubernetes cluster. The command has a simple and unique syntax to manage everything :

```
kubectl [command] [TYPE] [NAME] [flags]
```

- *command* : specifies the operation that you want to perform on one or more resources (create, get, describe, delete)
- *type* : specifies the resource type. Resource types are case-insensitive and you can specify the singular, plural, or abbreviated forms
- *name* : specifies the name of the resource. Names are case-sensitive. If the name is omitted, details for all resources are displayed
- *flags* : specifies optional flags.

Useful basic commands

Create

Create a resource from a file or from stdin.

```
# Create a pod using the data in pod.json.
```

```
kubectl create -f ./pod.json
```

```
# Create a pod based on the JSON passed into stdin.
```

```
cat pod.json | kubectl create -f -
```

```
# Edit the data in docker-registry.yaml in JSON using the v1 API format then create
```

```
kubectl create -f docker-registry.yaml --edit --output-version=v1 -o json
```

```
# Create all the resources available in the folder
```

```
kubectl create -f <folder_name>
```

Delete

Delete resources by filenames, stdin, resources and names, or by resources and label selector.

```
# Delete a pod using the type and name specified in pod.json.
kubectl delete -f ./pod.json

# Delete a pod based on the type and name in the JSON passed into stdin.
cat pod.json | kubectl delete -f -

# Delete pods and services with same names "baz" and "foo"
kubectl delete pod,service baz foo

# Delete pods and services with label name=myLabel.
kubectl delete pods,services -l name=myLabel

# Delete a pod with minimal delay
kubectl delete pod foo --now

# Force delete a pod on a dead node
kubectl delete pod foo --grace-period=0 --force

# Delete all pods
kubectl delete pods --all

#Delete all resources available in the folder.
kubectl delete -f <folder_name>
```

Edit

Edit a resource from the default editor.

```
# Edit the service named 'docker-registry':
kubectl edit svc/docker-registry

# Use an alternative editor
KUBE_EDITOR="nano" kubectl edit svc/docker-registry

# Edit the job 'myjob' in JSON using the v1 API format:
kubectl edit job.v1.batch/myjob -o json

# Edit the deployment 'mydeployment' in YAML and save the modified config in its ann
kubectl edit deployment/mydeployment -o yaml --save-config
```

Expose

Expose a resource as a new Kubernetes service.

```
# Create a service for a replicated nginx, which serves on port 80 and connects to t
kubectl expose rc nginx --port=80 --target-port=8000

# Create a service for a replication controller identified by type and name specifie
kubectl expose -f nginx-controller.yaml --port=80 --target-port=8000

# Create a service for a pod valid-pod, which serves on port 444 with the name "fron
kubectl expose pod valid-pod --port=444 --name=frontend

# Create a second service based on the above service, exposing the container port 84
kubectl expose service nginx --port=443 --target-port=8443 --name=nginx-https

# Create a service for a replicated streaming application on port 4100 balancing UDP
kubectl expose rc streamer --port=4100 --protocol=udp --name=video-stream

# Create a service for a replicated nginx using replica set, which serves on port 80
kubectl expose rs nginx --port=80 --target-port=8000

# Create a service for an nginx deployment, which serves on port 80 and connects to
kubectl expose deployment nginx --port=80 --target-port=8000

# Access Pod without exposing as service using kubectl --raw
kubectl get pod <pod_name> -o yaml|grep selfLink
kubectl get --raw <selfLink>:port/proxy/<filename>
```

Get

Display one or many resources.

```
# List all pods.
kubectl get pods

# List all pods in ps output format with more information (such as node name).
kubectl get pods -o wide

# List a single replication controller with specified NAME in ps output format.
kubectl get replicationcontroller web

# List a single pod in JSON output format.
kubectl get -o json pod <pod-name>

# List a pod identified by type and name specified in "pod.yaml" in JSON output form
kubectl get -f pod.yaml -o json

# Return only the phase value of the specified pod.
kubectl get -o template pod/<pod-name> --template=
```

```
# List all replication controllers and services together in ps output format.
kubectl get rc,services

# List one or more resources by their type and names.
kubectl get rc/web service/frontend pods/<pod-name>

# List all resources with different types.
kubectl get all
```

Run

Create and run a particular image, possibly replicated.

```
# Start a single instance of nginx.
kubectl run nginx --image=nginx

# Start a single instance of hazelcast and let the container expose port 5701 .
kubectl run hazelcast --image=hazelcast --port=5701

# Start a single instance of hazelcast and set environment variables "DNS_DOMAIN=cluster" and "POD_NAMESPACE=default"
kubectl run hazelcast --image=hazelcast --env="DNS_DOMAIN=cluster" --env="POD_NAMESPACE=default"

# Start a single instance of hazelcast and set labels "app=hazelcast" and "env=prod"
kubectl run hazelcast --image=nginx --labels="app=hazelcast,env=prod"

# Start a replicated instance of nginx.
kubectl run nginx --image=nginx --replicas=5

# Dry run. Print the corresponding API objects without creating them.
kubectl run nginx --image=nginx --dry-run

# Start a single instance of nginx, but overload the spec of the deployment with a pod
kubectl run nginx --image=nginx --overrides='{ "apiVersion": "v1", "spec": { ... } }'

# Start a pod of busybox and keep it in the foreground, don't restart it if it exits
kubectl run -i -t busybox --image=busybox --restart=Never

# Start the nginx container using the default command, but use custom arguments (arguments are passed to the container)
kubectl run nginx --image=nginx -- <arg1> <arg2> ... <argN>

# Start the nginx container using a different command and custom arguments.
kubectl run nginx --image=nginx --command -- <cmd> <arg1> ... <argN>

# Start the perl container to compute pi to 2000 places and print it out.
kubectl run pi --image=perl --restart=OnFailure -- perl -Mbignum=bpi -wle 'print bpi'
```

```
# Start the cron job to compute  $\pi$  to 2000 places and print it out every 5 minutes.  
kubectl run pi --schedule="0/5 * * * ?" --image=perl --restart=OnFailure -- perl -Mb
```

Set

Configure application resources.

```
# Update deployment 'registry' with a new environment variable  
kubectl set env deployment/registry STORAGE_DIR=/local  
  
# List the environment variables defined on a deployments 'sample-build'  
kubectl set env deployment/sample-build --list  
  
# List the environment variables defined on all pods  
kubectl set env pods --all --list  
  
# Output modified deployment in YAML, and does not alter the object on the server  
kubectl set env deployment/sample-build STORAGE_DIR=/data -o yaml  
  
# Update all containers in all replication controllers in the project to have ENV=prod  
kubectl set env rc --all ENV=prod  
  
# Import environment from a secret  
kubectl set env --from=secret/mysecret deployment/myapp  
  
# Import environment from a config map with a prefix  
kubectl set env --from=configmap/myconfigmap --prefix=MYSQL_ deployment/myapp  
  
# Remove the environment variable ENV from container 'c1' in all deployment configs  
kubectl set env deployments --all --containers="c1" ENV-  
  
# Remove the environment variable ENV from a deployment definition on disk and  
# update the deployment config on the server  
kubectl set env -f deploy.json ENV-  
  
# Set some of the local shell environment into a deployment config on the server  
env | grep RAILS_ | kubectl set env -e - deployment/registry  
  
# Set a deployment's nginx container image to 'nginx:1.9.1', and its busybox container  
kubectl set image deployment/nginx busybox=busybox nginx=nginx:1.9.1  
  
# Update all deployments' and rc's nginx container's image to 'nginx:1.9.1'  
kubectl set image deployments,rc nginx=nginx:1.9.1 --all  
  
# Update image of all containers of daemonset abc to 'nginx:1.9.1'
```

```
kubectl set image daemonset abc *=nginx:1.9.1

# Print result (in yaml format) of updating nginx container image from local file, w
kubectl set image -f path/to/file.yaml nginx=nginx:1.9.1 --local -o yaml

# Set a deployments nginx container cpu limits to "200m" and memory to "512Mi"
kubectl set resources deployment nginx -c=nginx --limits=cpu=200m,memory=512Mi

# Set the resource request and limits for all containers in nginx
kubectl set resources deployment nginx --limits=cpu=200m,memory=512Mi --requests=cpu

# Remove the resource requests for resources on containers in nginx
kubectl set resources deployment nginx --limits=cpu=0,memory=0 --requests=cpu=0,memo

# Print the result (in yaml format) of updating nginx container limits from a local,
kubectl set resources -f path/to/file.yaml --limits=cpu=200m,memory=512Mi --local -o

# Set Deployment nginx-deployment's ServiceAccount to serviceaccount1
kubectl set serviceaccount deployment nginx-deployment serviceaccount1

# Print the result (in yaml format) of updated nginx deployment with serviceaccount
kubectl set sa -f nginx-deployment.yaml serviceaccount1 --local --dry-run -o yaml
```

Useful deploy commands

Autoscale

Creates an autoscaler that automatically chooses and sets the number of pods that run in a kubernetes cluste

```
# Auto scale a deployment "foo", with the number of pods between 2 and 10, no target
kubectl autoscale deployment foo --min=2 --max=10

# Auto scale a replication controller "foo", with the number of pods between 1 and 5
kubectl autoscale rc foo --max=5 --cpu-percent=80
```

Rollout

Manage the rollout of a resource.

```
# Rollback to the previous deployment
kubectl rollout undo deployment/abc
```



```
# Check the rollout status of a daemonset
kubectl rollout status daemonset/foo

# View the rollout history of a deployment
kubectl rollout history deployment/abc

# View the details of daemonset revision 3
kubectl rollout history daemonset/abc --revision=3

# Mark the nginx deployment as paused. Any current state of
# the deployment will continue its function, new updates to the deployment will not
# have an effect as long as the deployment is paused.
kubectl rollout pause deployment/nginx

# Resume an already paused deployment
kubectl rollout resume deployment/nginx

# Watch the rollout status of a deployment
kubectl rollout status deployment/nginx

# Rollback to the previous deployment
kubectl rollout undo deployment/abc

# Rollback to daemonset revision 3
kubectl rollout undo daemonset/abc --to-revision=3

# Rollback to the previous deployment with dry-run
kubectl rollout undo --dry-run=true deployment/abc
```

Scale

Set a new size for a Deployment, ReplicaSet, Replication Controller, or StatefulSet.

```
# Scale a replicaset named 'foo' to 3.
kubectl scale --replicas=3 rs/foo

# Scale a resource identified by type and name specified in "foo.yaml" to 3.
kubectl scale --replicas=3 -f foo.yaml

# If the deployment named mysql's current size is 2, scale mysql to 3.
kubectl scale --current-replicas=2 --replicas=3 deployment/mysql

# Scale multiple replication controllers.
kubectl scale --replicas=5 rc/foo rc/bar rc/baz
```

```
# Scale statefulset named 'web' to 3.
kubectl scale --replicas=3 statefulset/web
```

Useful cluster management commands

Cluster-info

Display addresses of the master and services with label `kubernetes.io/cluster-service=true` To further debug and diagnose cluster problems, use `'kubectl cluster-info dump'`.

```
# Print the address of the master and cluster services
kubectl cluster-info
```

Cordon / Uncordon

Mark node as (un)schedulable.

```
# Mark node "foo" as unschedulable.
kubectl cordon foo

# Mark node "foo" as schedulable.
$ kubectl uncordon foo
```

Drain

Drain node in preparation for maintenance.

```
# Drain node "foo", even if there are pods not managed by a ReplicationController, R
$ kubectl drain foo --force

# As above, but abort if there are pods not managed by a ReplicationController, Repl
$ kubectl drain foo --grace-period=90

#Drain node by ignoring Deamonsets
kubectl drain <node_name> --ignore-daemonsets
```

Taint

Update the taints on one or more nodes.

```
# Update node 'foo' with a taint with key 'dedicated' and value 'special-user' and e
# If a taint with that key and effect already exists, its value is replaced as speci
kubectl taint nodes foo dedicated=special-user:NoSchedule

# Remove from node 'foo' the taint with key 'dedicated' and effect 'NoSchedule' if o
kubectl taint nodes foo dedicated:NoSchedule-

# Remove from node 'foo' all the taints with key 'dedicated'
kubectl taint nodes foo dedicated-

# Add a taint with key 'dedicated' on nodes having label myLabel=X
kubectl taint node -l myLabel=X dedicated=foo:PreferNoSchedule
```

Top

Display Resource (CPU/Memory/Storage) usage.

```
# Show metrics for all nodes
kubectl top node

# Show metrics for a given node
kubectl top node NODE_NAME

# Show metrics for all pods in the default namespace
kubectl top pod

# Show metrics for all pods in the given namespace
kubectl top pod --namespace=NAMESPACE

# Show metrics for a given pod and its containers
kubectl top pod POD_NAME --containers

# Show metrics for the pods defined by label name=myLabel
kubectl top pod -l name=myLabel
```

Useful troubleshooting and debugging commands

Describe

Show details of a specific resource or group of resources.

```
# Describe a node
kubectl describe nodes kubernetes-node-emt8.c.myproject.internal

# Describe a pod
kubectl describe pods/<pod-name>

# Describe a pod identified by type and name in "pod.json"
kubectl describe -f pod.json

# Describe all pods
kubectl describe pods

# Describe pods by label name=myLabel
kubectl describe po -l name=myLabel

# Describe all pods managed by the 'frontend' replication controller (rc-created pod
# get the name of the rc as a prefix in the pod the name).
kubectl describe pods frontend
```

Exec

Execute a command in a container.

```
# Get output from running 'date' from pod 123456-7890, using the first container by
kubectl exec 123456-7890 date

# Get output from running 'date' in ruby-container from pod 123456-7890
kubectl exec 123456-7890 -c ruby-container date

# Switch to raw terminal mode, sends stdin to 'bash' in ruby-container from pod 1234
# and sends stdout/stderr from 'bash' back to the client
kubectl exec 123456-7890 -c ruby-container -i -t -- bash -il

# List contents of /usr from the first container of pod 123456-7890 and sort by modi
# If the command you want to execute in the pod has any flags in common (e.g. -i),
# you must use two dashes (--) to separate your command's flags/arguments.
# Also note, do not surround your command and its flags/arguments with quotes
# unless that is how you would execute it normally (i.e., do ls -t /usr, not "ls -t
kubectl exec 123456-7890 -i -t -- ls -t /usr
```

Logs

Print the logs for a container in a pod or specified resource. If the pod has only one container, the container name is optional.

```
# Return snapshot logs from pod nginx with only one container
kubectl logs nginx

# Return snapshot logs for the pods defined by label app=nginx
kubectl logs -lapp=nginx

# Return snapshot of previous terminated ruby container logs from pod web-1
kubectl logs -p -c ruby web-1

# Begin streaming the logs of the ruby container in pod web-1
kubectl logs -f -c ruby web-1

# Display only the most recent 20 lines of output in pod nginx
kubectl logs --tail=20 nginx

# Show all logs from pod nginx written in the last hour
kubectl logs --since=1h nginx

# Return snapshot logs from first container of a job named hello
kubectl logs job/hello

# Return snapshot logs from container nginx-1 of a deployment named nginx
kubectl logs deployment/nginx -c nginx-1
```

Proxy

Creates a proxy server or application-level gateway between localhost and the Kubernetes API Server. It also allows serving static content over specified HTTP path. All incoming data enters through one port and gets forwarded to the remote kubernetes API Server port, except for the path matching the static content path.

```
# To proxy all of the kubernetes api and nothing else, use:
$ kubectl proxy --api-prefix=/

# To proxy only part of the kubernetes api and also some static files:
$ kubectl proxy --www=/my/files --www-prefix=/static/ --api-prefix=/api/
# The above lets you 'curl localhost:8001/api/v1/pods'.

# To proxy the entire kubernetes api at a different root, use:
$ kubectl proxy --api-prefix=/custom/
# The above lets you 'curl localhost:8001/custom/api/v1/pods'

# Run a proxy to kubernetes apiserver on port 8011, serving static content from ./lo
kubectl proxy --port=8011 --www=./local/www/
```

```
# Run a proxy to kubernetes apiserver on an arbitrary local port.  
# The chosen port for the server will be output to stdout.  
kubectl proxy --port=0
```

Useful advanced commands

Apply

Apply a configuration to a resource by filename or stdin. The resource name must be specified. This resource will be created if it doesn't exist yet. To use 'apply', always create the resource initially with either 'apply' or 'create --save-config'.

```
# Apply the configuration in pod.json to a pod.  
kubectl apply -f ./pod.json  
  
# Apply the JSON passed into stdin to a pod.  
cat pod.json | kubectl apply -f -  
  
# Note: --prune is still in Alpha  
# Apply the configuration in manifest.yaml that matches label app=nginx and delete a  
kubectl apply --prune -f manifest.yaml -l app=nginx  
  
# Apply the configuration in manifest.yaml and delete all the other configmaps that  
kubectl apply --prune -f manifest.yaml --all --prune-whitelist=core/v1/ConfigMap
```

Useful settings commands

label

Update the labels on a resource.

```
# Update pod 'foo' with the label 'unhealthy' and the value 'true'.  
kubectl label pods foo unhealthy=true  
  
# Update pod 'foo' with the label 'status' and the value 'unhealthy', overwriting an  
kubectl label --overwrite pods foo status=unhealthy  
  
# Update all pods in the namespace  
kubectl label pods --all status=unhealthy  
  
# Update a pod identified by the type and name in "pod.json"
```

```
kubectl label -f pod.json status=unhealthy

# Update pod 'foo' only if the resource is unchanged from version 1.
kubectl label pods foo status=unhealthy --resource-version=1

# Update pod 'foo' by removing a label named 'bar' if it exists.
# Does not require the --overwrite flag.
kubectl label pods foo bar-
```

Useful other commands

Config

Modify kubeconfig files using subcommands like “kubectl config set current-context my-context”.

```
# Display the current-context
kubectl config current-context

# Delete the minikube cluster
kubectl config delete-cluster minikube

# Delete the context for the minikube cluster
kubectl config delete-context minikube

# List the clusters kubectl knows about
kubectl config get-clusters

# List the context kubectl knows about
kubectl config get-contexts

# Rename the context 'old-name' to 'new-name' in your kubeconfig file
kubectl config rename-context old-name new-name

# Set only the server field on the e2e cluster entry without touching other values.
kubectl config set-cluster e2e --server=https://1.2.3.4

# Embed certificate authority data for the e2e cluster entry
kubectl config set-cluster e2e --certificate-authority=~/.kube/e2e/kubernetes.ca.crt

# Disable cert checking for the dev cluster entry
kubectl config set-cluster e2e --insecure-skip-tls-verify=true

# Set the user field on the gce context entry without touching other values
kubectl config set-context gce --user=cluster-admin
```

```
# Use the context for the minikube cluster
kubectl config use-context minikube
```

Version

Print the client and server version information for the current context.

```
# Print the client and server versions for the current context
kubectl version
```

External documentation


To go further in the management of Kubectl, please refer to these documentations :

- Official Kubernetes [overview of Kubectl](#) command line
- Official Kubernetes documentation to [install Kubectl](#) command line
- Official [Kubectl commands](#) details

Contributors

- [Bala](#) - balasundarammaster@gmail.com
- [Wikitops](#) - wikiposts5692@gmail.com
- [Vinayak Ladwa](#) - vinayakladwa@gmail.com

Did you find Dockerlabs useful? Vote for us.



Dockerlabs - Docker and Kubernetes

(dockerlabs.collabnix.com)

Submitted by Sangam | 6 months ago | 3.7k+ views

Free

> [hackr.io](#)

Tweets by [collabnix](#)

dockerlabs is maintained by **collabnix**.

This page was generated by [GitHub Pages](#).